

HW0_Final.R

James

2019-04-13

```
setwd("~/GitHub/MMSS_311_2")

# Q1a) A vector with the numbers 1-5 in order
a <- c(1,2,3,4,5)

# Q1b) A scalar named Mindy that takes the value 12
Mindy <- 12

# Q1c) A 2x3 matrix with the numbers 1-6 in order by rows
b <- c(1,2,3,4,5,6)
matrix(b,2,3,TRUE)
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    3
## [2,]    4    5    6
```

```
# Q1d) A 2x3 matrix with the numbers 1-6 in order by columns
matrix(b,2,3)
```

```
##      [,1] [,2] [,3]
## [1,]    1    3    5
## [2,]    2    4    6
```

```
# Q1e) A 10x10 matrix of 1's
matrix(1,10,10)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
## [1,]    1    1    1    1    1    1    1    1    1    1
## [2,]    1    1    1    1    1    1    1    1    1    1
## [3,]    1    1    1    1    1    1    1    1    1    1
## [4,]    1    1    1    1    1    1    1    1    1    1
## [5,]    1    1    1    1    1    1    1    1    1    1
## [6,]    1    1    1    1    1    1    1    1    1    1
## [7,]    1    1    1    1    1    1    1    1    1    1
## [8,]    1    1    1    1    1    1    1    1    1    1
## [9,]    1    1    1    1    1    1    1    1    1    1
## [10,]   1    1    1    1    1    1    1    1    1    1
```

```
# Q1f) A vector consisting of the words THIS, IS, A, VECTOR (each word a separate element)
wordvec <- c("THIS", "IS", "A", "VECTOR")

# Q1g) A function that takes the sum of any three numbers
sum_of_three_numbers <- function(x,y,z) {
  x+y+z
}

# Q1h) A function that takes one number as input, returns "Yes" if the number is less than or equal to 10 and "No" if the number is greater than 10
check <- function(x) {
  if (x<=10) {
    result <- "Yes"
  }
  else if (x>10) {
    result <- "No"
  }
  return(result)
}
check(9)
```

```
## [1] "Yes"
```

```
# Q1i) Generate synthetic data by taking 1,000 draws from a normal distribution with a mean of 10 and a standard deviation of 1. Save these data to an object g.
g <- rnorm(1000,10,1)

# Q1j) Create a separate object called y with 1,000 draws from a normal distribution with a mean of 5 and a standard deviation of 0.5.
y <- rnorm(1000,5,0.5)

# Q1k) Generate a variable x with 1,000 values, where each value is a mean of 10 samples from g, with replacement. (Hint: use a for loop)
x = NULL
for(i in 1:1000) {
  x [i] <- mean(sample(g, 10, TRUE))
}

# Q1l) Estimate a simple bivariate regression y on x and print your results. What do your results show?
# The results show that the OLS estimator of the coefficient of x is 0.02633, which is very small. This shows that there is only a weak positive correlation between y and x.
reg <- lm(y ~ x)
print(reg)
```

```
##  
## Call:  
## lm(formula = y ~ x)  
##  
## Coefficients:  
## (Intercept)          x  
##      4.90030      0.01247
```

```
#Q2a Create an R script file that sets your working directory and loads the data.  
setwd("~/GitHub/MMSS_311_2")
```

```
pums_chicago <- read.csv("pums_chicago.csv")
```

```
#2b How many variables are there in the dataset?  
# There are 204 variables (from environment panel)
```

```
#2c What is the mean annual income, PINCP in this dataset?  
PINCP_mean <- mean(pums_chicago$PINCP, na.rm = TRUE)
```

```
#2d Create a new variable in the PUMS dataframe called PINCP_LOG that is equal to the log of annual income. Were NaN values produced? Why?  
#NaN values were produced because we cannot take log of 0, which is the value of some annual income observations.  
pums_chicago$PINCP_LOG <- log(pums_chicago$PINCP)
```

```
## Warning in log(pums_chicago$PINCP): NaNs produced
```

#2e Create a new variable GRAD.DUMMY that takes the value "grad" if the respondent has any post-high school education, and "no grad" otherwise. Use the SCHL variable.

```
pums_chicago$GRAD.DUMMY <- ifelse(pums_chicago$SCHL > 17, "grad", "no grad")
```

#2f Drop the variable SERIALNO from the dataset.

```
pums_chicago$SERIALNO <- NULL
```

#2g Save your new dataset to a csv file in the working directory.

```
write.csv(pums_chicago, 'editedPUMS_CHICAGO.csv')
```

#2h Use the variable ESR, create 5 new dataframes: under 16, employed, unemployed, in the armed forces, and not in the labor force.

```
under16 <- pums_chicago[pums_chicago$ESR == "NA", ]
employed <- pums_chicago[pums_chicago$ESR %in% c("1", "2"), ]
unemployed <- pums_chicago[pums_chicago$ESR %in% "3", ]
armedforces <- pums_chicago[pums_chicago$ESR %in% c("4", "5"), ]
notinlaborforce <- pums_chicago[pums_chicago$ESR %in% "6", ]
```

#2i Create a new dataframe that combines employed people and people in the armed forces.

```
employed_af <- pums_chicago[pums_chicago$ESR %in% c("1", "2", "4", "5"), ]
```

#2j In your new employed_af dataframe, keep only the variables AGEP, RAC1P, and PINCP_LOG

```
new_employed_af <- pums_chicago[c("AGEP", "RAC1P", "PINCP_LOG")]
```

#2ki Find the mean, median, and 80th percentile of travel time to work, JWMNP

```
summary(pums_chicago$JWMNP)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##      1.00   20.00   30.00   34.84   45.00  149.00  27668
```

```
quantile(pums_chicago$JWMNP, probs=0.8, na.rm=TRUE)
```

```
## 80%
```

```
## 45
```

#2kii Find the correlation between travel time to work JWMNP and annual wages WAGP

```
cor(pums_chicago$JWMNP, pums_chicago$WAGP, use="complete.obs")
```

```
## [1] -0.04205232
```

#2kiii Make a scatterplot of age and log income.

#2kiv Export this graph to your working directory in pdf format.

```
pdf("ageonLogincome.pdf")
plot(pums_chicago$AGEP, pums_chicago$PINCP_LOG, main="age on log income")
dev.off()
```

```
## png
```

```
## 2
```

```
#2kv Create a crosstab of employment status ESR by race RAC1P  
install.packages("gmodels", repos = "http://cran.us.r-project.org")
```

```
## Installing package into 'C:/Users/James/Documents/R/win-library/3.5'  
## (as 'lib' is unspecified)
```

```
## package 'gmodels' successfully unpacked and MD5 sums checked  
##  
## The downloaded binary packages are in  
## C:\Users\James\AppData\Local\Temp\RtmpgvKp2C\downloaded_packages
```

```
library("gmodels")  
CrossTable(pums_chicago$ESR, pums_chicago$RAC1P)
```

```
##
##
##   Cell Contents
## |-----|
## |               N |
## | Chi-square contribution |
## |           N / Row Total |
## |           N / Col Total |
## |           N / Table Total |
## |-----|
##
##
## Total Observations in Table:  40348
##
##
##                | pums_chicago$RAC1P
## pums_chicago$ESR |           1 |           2 |           3 |           4 |           5 |           6 |
7 |           8 |           9 | Row Total |
## -----|-----|-----|-----|-----|-----|-----|
## -----|-----|-----|-----|-----|-----|-----|
##                1 |      12870 |      5786 |           36 |           0 |           24 |           1746 |
7 |      2502 |           521 |      23492 |
##                |      195.313 |      406.565 |      0.689 |      1.164 |      0.414 |      9.555 |
0.591 |      4.491 |      3.239 |           |
##                |      0.548 |      0.246 |      0.002 |      0.000 |      0.001 |      0.074 |
0.000 |      0.107 |      0.022 |      0.582 |
##                |      0.659 |      0.447 |      0.507 |      0.000 |      0.511 |      0.627 |
0.778 |      0.607 |      0.630 |           |
##                |      0.319 |      0.143 |      0.001 |      0.000 |      0.001 |      0.043 |
0.000 |      0.062 |      0.013 |           |
## -----|-----|-----|-----|-----|-----|-----|
## -----|-----|-----|-----|-----|-----|-----|
##                2 |           258 |           147 |           0 |           0 |           0 |           31 |
0 |           66 |           8 |           510 |
##                |           0.487 |           1.687 |      0.897 |      0.025 |      0.594 |      0.502 |
0.114 |      3.730 |           0.576 |           |
##                |           0.506 |           0.288 |      0.000 |      0.000 |      0.000 |      0.061 |
0.000 |      0.129 |           0.016 |           0.013 |
##                |           0.013 |           0.011 |      0.000 |      0.000 |      0.000 |      0.011 |
0.000 |      0.016 |           0.010 |           |
##                |           0.006 |           0.004 |      0.000 |      0.000 |      0.000 |      0.001 |
0.000 |      0.002 |           0.000 |           |
## -----|-----|-----|-----|-----|-----|-----|
## -----|-----|-----|-----|-----|-----|-----|
##                3 |           794 |           1473 |           2 |           0 |           4 |           109 |
0 |           268 |           57 |           2707 |
##                |           204.029 |           420.880 |      1.603 |      0.134 |      0.227 |      32.435 |
0.604 |      0.252 |           0.041 |           |
##                |           0.293 |           0.544 |      0.001 |      0.000 |      0.001 |      0.040 |
0.000 |      0.099 |           0.021 |           0.067 |
##                |           0.041 |           0.114 |      0.028 |      0.000 |      0.085 |      0.039 |
0.000 |      0.065 |           0.069 |           |
##                |           0.020 |           0.037 |      0.000 |      0.000 |      0.000 |      0.003 |
```

```

0.000 |    0.007 |    0.001 |    |
## -----|-----|-----|-----|-----|-----|-----|-----|
-----|-----|-----|-----|
##          4 |          4 |          5 |          0 |          0 |          0 |          0 |
1 |          0 |          1 |         11 |
##          |          0.331 |          0.613 |          0.019 |          0.001 |          0.013 |          0.759 |
405.558 |          1.123 |          2.661 |          |
##          |          0.364 |          0.455 |          0.000 |          0.000 |          0.000 |          0.000 |
0.091 |          0.000 |          0.091 |          0.000 |
##          |          0.000 |          0.000 |          0.000 |          0.000 |          0.000 |          0.000 |
0.111 |          0.000 |          0.001 |          |
##          |          0.000 |          0.000 |          0.000 |          0.000 |          0.000 |          0.000 |
0.000 |          0.000 |          0.000 |          |
## -----|-----|-----|-----|-----|-----|-----|-----|
-----|-----|-----|-----|
##          6 |          5618 |          5533 |          33 |          2 |          19 |          899 |
1 |          1283 |          240 |          13628 |
##          |          146.443 |          308.317 |          3.392 |          2.597 |          0.615 |          1.846 |
1.369 |          8.421 |          5.537 |          |
##          |          0.412 |          0.406 |          0.002 |          0.000 |          0.001 |          0.066 |
0.000 |          0.094 |          0.018 |          0.338 |
##          |          0.287 |          0.427 |          0.465 |          1.000 |          0.404 |          0.323 |
0.111 |          0.311 |          0.290 |          |
##          |          0.139 |          0.137 |          0.001 |          0.000 |          0.000 |          0.022 |
0.000 |          0.032 |          0.006 |          |
## -----|-----|-----|-----|-----|-----|-----|-----|
-----|-----|-----|-----|
## Column Total |          19544 |          12944 |          71 |          2 |          47 |          2785 |
9 |          4119 |          827 |          40348 |
##          |          0.484 |          0.321 |          0.002 |          0.000 |          0.001 |          0.069 |
0.000 |          0.102 |          0.020 |          |
## -----|-----|-----|-----|-----|-----|-----|-----|
-----|-----|-----|-----|
##
##

```

```
#2kvi Estimate a linear regression of annual wages WAGP on hours worked per week WKHP
wagp_on_wkhp <- lm(WAGP ~ WKHP, pums_chicago)

#2kvii Plot the residuals from this regression against the fitted values. What does this show?
#This shows that residuals tend to decrease as the fitted values increase. Furthermore, we can observe that there are only a small number of large deviations for any given fitted value.
wagp_on_wkhp_res <- resid(wagp_on_wkhp)
wagp_on_wkhp_fitted <- fitted(wagp_on_wkhp)
plot(wagp_on_wkhp_fitted, wagp_on_wkhp_res,
     ylab = "Residuals", xlab="Fitted",
     main = "Residuals on Fitted")

#2li Estimate a linear regression of miles per gallon on weight
data(mtcars)
car_lm <- lm(mpg ~ wt, mtcars)

#2lii Estimate this regression separately for manual versus automatic transmission
autoData <- mtcars[mtcars$am == "0",]
manualData <- mtcars[mtcars$am == "1",]
autocar_lm <- lm(mpg ~ wt, mtcars)
manualcar_lm <- lm(mpg ~ wt, mtcars)


#2liii Estimate a regression of miles per gallon on the log of horsepower.
mtcars$log.hp <- log(mtcars$hp)
mpg_on_lg.hp <- lm(mpg ~ log.hp, mtcars)

#2mi Make a scatterplot of weight against miles per gallon.
#2mii Color the points in your graph according to the transmission of the vehicle.
#2miii Change the shape of the points to correspond to the number of forward gears in the vehicle.
#2miv Change the x and y labels on the plot to make full words.
#2mv Change the background of the plot so that the panel background is not gray.
install.packages("ggplot2", repos = "http://cran.us.r-project.org")
```

```
## Installing package into 'C:/Users/James/Documents/R/win-library/3.5'
## (as 'lib' is unspecified)
```

```
## package 'ggplot2' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## C:\Users\James\AppData\Local\Temp\RtmpgvKp2C\downloaded_packages
```

```
library(ggplot2)
```

```
ggplot(mtcars)+ geom_point(mapping = aes(x = mpg, y = wt,
                                         color = mtcars$am,
                                         shape = mtcars$gear)) + scale_shape_identity() +
labs(title = "Weight on Miles per Gallon", x = "Miles per Gallon", y = "Weight") +
theme(panel.background = element_rect(fill = "lightyellow"))
```

